

Dataset 2 kpca+lda

December 30, 2022

```
[ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.utils import resample
from sklearn.preprocessing import scale
from sklearn.decomposition import PCA
from sklearn import metrics

from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.model_selection import validation_curve
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
import matplotlib.pyplot as plt
import matplotlib.colors as colors
import os
```

```
[ ]: csv0= pd.read_csv("0.csv")
csv1= pd.read_csv("1.csv")
csv2= pd.read_csv("2.csv")
csv3= pd.read_csv("3.csv")
```

```
[ ]: allFiles=['0.csv', '1.csv', '2.csv', '3.csv']
list = []
for file in allFiles:
    read = pd.read_csv(file, header = None)
    list.append(read)
data = pd.concat(list)
```

```
[ ]: data.info
```

```
[ ]: <bound method DataFrame.info of          0      1      2      3      4      5      6
7      8      9  ...    55  \
0      26.0   4.0   5.0   8.0   -1.0  -13.0 -109.0 -66.0   -9.0   2.0  ... -28.0
1     -47.0  -6.0  -5.0  -7.0   13.0   -1.0   35.0 -10.0   10.0  -4.0  ... -25.0
2     -19.0  -8.0  -8.0  -8.0  -21.0   -6.0  -79.0  12.0    0.0   5.0  ... -83.0
3         2.0   3.0   0.0   2.0    0.0   22.0  106.0 -14.0 -16.0  -2.0  ... -38.0
```

```

4      6.0  0.0  0.0 -2.0 -14.0  10.0 -51.0   5.0   7.0  0.0 ... 38.0
...
2917  -3.0 -1.0 -1.0 -1.0 -28.0  20.0   5.0   0.0 -5.0  0.0 ... -3.0
2918 -13.0 -5.0 -4.0 -3.0  -4.0 -24.0 -10.0  -8.0  20.0  9.0 ...  6.0
2919  -1.0 -3.0 -1.0  1.0  30.0  38.0  -1.0  36.0 -10.0  1.0 ... 14.0
2920   1.0  4.0  4.0  5.0   9.0 -10.0   4.0   1.0 -2.0 -1.0 ... -16.0
2921  -2.0  4.0  2.0 -4.0  12.0   3.0  -2.0   9.0 -8.0 -2.0 ...  2.0

```

```

      56  57  58  59  60  61  62  63  64
0      61.0  4.0  8.0  5.0  4.0 -7.0 -59.0 16.0  0
1      47.0  6.0  6.0  5.0 13.0 21.0 111.0 15.0  0
2       7.0  7.0  1.0 -8.0  7.0 21.0 114.0 48.0  0
3     -11.0  4.0  7.0 11.0 33.0 39.0 119.0 43.0  0
4    -35.0 -8.0  2.0  6.0 -13.0 -24.0 -112.0 -69.0  0
...
2917   1.0  4.0  3.0  4.0 -51.0 -49.0   5.0 -9.0  3
2918  -3.0 -3.0 -3.0 -5.0 -4.0 -45.0 -12.0 -15.0  3
2919  -8.0 -4.0 -4.0 -4.0 -21.0 -29.0  -5.0  0.0  3
2920  -3.0  0.0 -3.0 -5.0 -36.0 -90.0   3.0  5.0  3
2921   1.0  0.0 -1.0 -2.0 -30.0  64.0  11.0  5.0  3

```

```
[11678 rows x 65 columns]>
```

```
[ ]: data.isnull().sum().head
```

```

[ ]: <bound method NDFrame.head of 0      0
      1      0
      2      0
      3      0
      4      0
      ..
     60      0
     61      0
     62      0
     63      0
     64      0
Length: 65, dtype: int64>

```

```
[ ]: order= data[64].unique()
      print(order)
```

```
[0 1 2 3]
```

```
[ ]: data.head()
```

```

[ ]:      0      1      2      3      4      5      6      7      8      9      ...      55      56  \
0  26.0  4.0  5.0  8.0  -1.0 -13.0 -109.0 -66.0  -9.0  2.0 ... -28.0  61.0
1 -47.0 -6.0 -5.0 -7.0  13.0  -1.0   35.0 -10.0  10.0 -4.0 ... -25.0  47.0

```

```

2 -19.0 -8.0 -8.0 -8.0 -21.0 -6.0 -79.0 12.0 0.0 5.0 ... -83.0 7.0
3 2.0 3.0 0.0 2.0 0.0 22.0 106.0 -14.0 -16.0 -2.0 ... -38.0 -11.0
4 6.0 0.0 0.0 -2.0 -14.0 10.0 -51.0 5.0 7.0 0.0 ... 38.0 -35.0

```

```

      57  58  59  60  61  62  63  64
0  4.0  8.0  5.0  4.0 -7.0 -59.0 16.0  0
1  6.0  6.0  5.0 13.0 21.0 111.0 15.0  0
2  7.0  1.0 -8.0  7.0 21.0 114.0 48.0  0
3  4.0  7.0 11.0 33.0 39.0 119.0 43.0  0
4 -8.0  2.0  6.0 -13.0 -24.0 -112.0 -69.0  0

```

[5 rows x 65 columns]

```
[ ]: y = data[64]
     X= data.drop(columns=64)
```

```
[ ]: data[64].value_counts().sort_values(ascending=False)
```

```

[ ]: 2    2943
     3    2922
     0    2910
     1    2903
     Name: 64, dtype: int64

```

```
[ ]: X.head
```

```

[ ]: <bound method NDFrame.head of          0    1    2    3    4    5    6    7
8    9  ...  54  \
0    26.0  4.0  5.0  8.0 -1.0 -13.0 -109.0 -66.0 -9.0  2.0 ...  21.0
1   -47.0 -6.0 -5.0 -7.0 13.0 -1.0  35.0 -10.0 10.0 -4.0 ... -105.0
2   -19.0 -8.0 -8.0 -8.0 -21.0 -6.0 -79.0 12.0  0.0  5.0 ... -128.0
3     2.0  3.0  0.0  2.0  0.0 22.0 106.0 -14.0 -16.0 -2.0 ...  -54.0
4     6.0  0.0  0.0 -2.0 -14.0 10.0 -51.0  5.0  7.0  0.0 ...   60.0
...
2917  -3.0 -1.0 -1.0 -1.0 -28.0 20.0    5.0  0.0 -5.0  0.0 ...  -3.0
2918 -13.0 -5.0 -4.0 -3.0 -4.0 -24.0 -10.0 -8.0 20.0  9.0 ...   5.0
2919  -1.0 -3.0 -1.0  1.0 30.0 38.0  -1.0 36.0 -10.0  1.0 ...  12.0
2920   1.0  4.0  4.0  5.0  9.0 -10.0   4.0  1.0 -2.0 -1.0 ...  -2.0
2921  -2.0  4.0  2.0 -4.0 12.0  3.0  -2.0  9.0 -8.0 -2.0 ... -10.0

```

```

      55  56  57  58  59  60  61  62  63
0  -28.0 61.0  4.0  8.0  5.0  4.0 -7.0 -59.0 16.0
1  -25.0 47.0  6.0  6.0  5.0 13.0 21.0 111.0 15.0
2  -83.0  7.0  7.0  1.0 -8.0  7.0 21.0 114.0 48.0
3  -38.0 -11.0  4.0  7.0 11.0 33.0 39.0 119.0 43.0
4   38.0 -35.0 -8.0  2.0  6.0 -13.0 -24.0 -112.0 -69.0
...

```

```

2917  -3.0   1.0   4.0   3.0   4.0 -51.0 -49.0    5.0  -9.0
2918   6.0  -3.0  -3.0  -3.0  -5.0  -4.0 -45.0  -12.0 -15.0
2919  14.0  -8.0  -4.0  -4.0  -4.0 -21.0 -29.0   -5.0   0.0
2920 -16.0  -3.0   0.0  -3.0  -5.0 -36.0 -90.0    3.0   5.0
2921   2.0   1.0   0.0  -1.0  -2.0 -30.0  64.0   11.0   5.0

```

[11678 rows x 64 columns]>

```
[ ]: X.describe()
```

```

[ ]:
count      0      1      2      3      4  \
mean      -0.520380  -0.726837  -0.739082  -0.729748  -0.159103
std       18.566709  11.766878   4.989944   7.441675  17.850402
min      -116.000000 -104.000000 -33.000000 -75.000000 -121.000000
25%       -9.000000  -4.000000  -3.000000  -4.000000 -10.000000
50%       -1.000000  -1.000000  -1.000000  -1.000000   0.000000
75%        7.000000   3.000000   2.000000   3.000000  10.000000
max       111.000000  90.000000  34.000000  55.000000  92.000000

count      5      6      7      8      9  \
mean      -0.554890  -1.272649  -0.661843  -0.665953  -0.654222
std       25.809528  25.089972  15.408896  18.123854  11.841260
min      -122.000000 -128.000000 -128.000000 -110.000000 -128.000000
25%      -15.000000  -6.000000  -8.000000  -9.000000  -4.000000
50%       -1.000000  -1.000000  -1.000000  -1.000000  -1.000000
75%       13.000000   4.000000   6.000000   6.000000   3.000000
max       127.000000  127.000000  126.000000  127.000000  106.000000

count  ...      54      55      56      57  \
mean  ...    -1.622452  -0.932694  -0.836958  -0.740623
std   ...    25.293001  15.158993  18.204465  12.005206
min   ...   -128.000000 -128.000000 -116.000000 -128.000000
25%   ...    -6.000000  -8.000000  -9.000000  -4.000000
50%   ...    -1.000000  -1.000000  -1.000000  -1.000000
75%   ...     3.000000   6.000000   6.000000   3.000000
max   ...    127.000000  114.000000  127.000000  105.000000

count      58      59      60      61      62  \
mean      -0.768710  -0.705343  -0.146686  -0.374807  -1.449306
std       4.969758   7.384410  17.841479  25.551082  25.259736
min      -46.000000  -74.000000 -103.000000 -128.000000 -128.000000
25%       -3.000000  -4.000000 -10.000000 -14.000000  -6.000000
50%       -1.000000  -1.000000   0.000000  -1.000000  -1.000000

```

75%	2.000000	3.000000	10.000000	13.000000	3.000000
max	29.000000	51.000000	110.000000	127.000000	127.000000

```

count 11678.000000
mean   -0.609094
std     15.530091
min    -124.000000
25%    -8.000000
50%    -1.000000
75%     6.000000
max    127.000000

```

[8 rows x 64 columns]

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X,y,random_state=40)
```

```
[ ]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled_train = scaler.fit_transform(X_train)
X_scaled_test = scaler.transform(X_test)
X_scaled_train = pd.DataFrame(X_scaled_train)
X_scaled_test = pd.DataFrame(X_scaled_test)
```

```
[ ]: X_scaled_test.head()
```

```
[ ]:
      0      1      2      3      4      5      6  \
0 -0.897962  0.223606  2.356874 -0.158795  0.349809 -0.015352  2.128207
1  0.026488 -0.374961 -1.253112 -0.556220  1.577761  0.448405 -0.070962
2 -1.441756 -0.118432  0.952990  0.503581 -0.878143 -0.672341 -0.030977
3  0.298385 -0.118432  0.150771 -0.026320  0.461441  1.298627 -0.030977
4  0.515903 -0.032923 -2.255886 -0.821171 -0.319983 -0.749634 -0.230902

      7      8      9  ...      54      55      56      57  \
0 -0.413754  0.313513 -0.624990  ... -1.480377 -0.196189  0.214258 -0.185965
1  1.075626  0.368848 -0.027987  ...  0.055997 -0.196189  3.193393  2.106044
2 -0.478510  0.866868 -0.113273  ...  0.055997 -0.196189 -0.282265  0.153592
3 -0.219488 -0.239842 -0.198559  ...  0.055997 -0.065461 -0.613280 -0.525522
4 -0.413754  0.756197  2.104165  ...  0.016602  0.261359  0.324596 -3.411756

      58      59      60      61      62      63
0  0.566839 -0.430394  1.484974  0.090064  2.208949  0.618095
1  2.575162  0.238175  0.061796  1.114511  0.298683  1.899496
2  0.366007  0.238175 -0.621330 -0.579766  0.219089  0.489955
3 -0.236491 -0.430394 -1.361383 -1.210194  0.020103 -0.342955
4 -0.236491  0.505602  0.346431  0.287073  0.219089 -0.150745
```

[5 rows x 64 columns]

```
[ ]: cov_matrix = X_scaled_train.cov()
eigen_vals, eigen_vecs = np.linalg.eigh(cov_matrix)
np.sum(eigen_vals)
```

```
[ ]: 64.00730843896312
```

```
[ ]: A = X_scaled_train.cov()
x = eigen_vecs[:, 0] # First eigenvector
LHS = np.dot(A, x) # Ax
print(LHS)
```

```
[-0.00071933 -0.00117752  0.00495437 -0.00958084  0.0111963  -0.00681954
 0.00117527  0.00129149 -0.00051435 -0.00161002  0.00965168 -0.01595066
 0.02074834 -0.01339627  0.00195487  0.0013108  0.00053154 -0.00184558
 0.01201653 -0.02024228  0.02896049 -0.01930371  0.00205734  0.00175303
 0.00042843 -0.00189969  0.01301459 -0.0214747  0.03104015 -0.02091013
 0.00225011  0.00144918  0.00082467 -0.00304437  0.01338423 -0.02178359
 0.03055792 -0.0214959  0.00311255  0.00147683  0.00170481 -0.00344283
 0.01287558 -0.01990288  0.02717143 -0.01865322  0.00176285  0.00090025
 0.00224643 -0.00370848  0.00940975 -0.01536319  0.01964771 -0.01506809
 0.00093975  0.00097418  0.00103895 -0.00270651  0.00552519 -0.00867789
 0.01036569 -0.00856539  0.00067139 -0.00021925]
```

```
[ ]: lambda_eig = eigen_vals[0] # First eigenvalue
x = eigen_vecs[:, 0] # First eigenvector
RHS = lambda_eig * x # x
print(RHS)
```

```
[-0.00071933 -0.00117752  0.00495437 -0.00958084  0.0111963  -0.00681954
 0.00117527  0.00129149 -0.00051435 -0.00161002  0.00965168 -0.01595066
 0.02074834 -0.01339627  0.00195487  0.0013108  0.00053154 -0.00184558
 0.01201653 -0.02024228  0.02896049 -0.01930371  0.00205734  0.00175303
 0.00042843 -0.00189969  0.01301459 -0.0214747  0.03104015 -0.02091013
 0.00225011  0.00144918  0.00082467 -0.00304437  0.01338423 -0.02178359
 0.03055792 -0.0214959  0.00311255  0.00147683  0.00170481 -0.00344283
 0.01287558 -0.01990288  0.02717143 -0.01865322  0.00176285  0.00090025
 0.00224643 -0.00370848  0.00940975 -0.01536319  0.01964771 -0.01506809
 0.00093975  0.00097418  0.00103895 -0.00270651  0.00552519 -0.00867789
 0.01036569 -0.00856539  0.00067139 -0.00021925]
```

```
[ ]: explained_variance=[]

for i in sorted(eigen_vals,reverse=True):
    variance = (i / np.sum(eigen_vals))
    explained_variance.append(variance)
```

```
cumulative_variance_ratio = np.cumsum(explained_variance)*100
print(cumulative_variance_ratio)
```

```
[ 4.70458409  9.05851929 13.01124542 16.91114022 20.66601573
 24.11699848 27.50135433 30.7106687  33.86599837 36.87929532
 39.78712524 42.46965224 45.08342903 47.65124562 50.15999691
 52.58691677 54.94842403 57.04380323 59.04692089 61.04374454
 62.97559477 64.837373  66.67903344 68.26566973 69.81758623
 71.33211173 72.74349853 74.10466444 75.43097003 76.65159236
 77.83110067 78.99758871 80.14028338 81.25562719 82.33970153
 83.40460512 84.42191333 85.41285184 86.37641944 87.33211569
 88.26160584 89.09448658 89.88205271 90.61091588 91.31954074
 92.0086025  92.68720904 93.35663384 93.98099591 94.5973542
 95.19775037 95.76188013 96.29341372 96.80876446 97.30090091
 97.7533043  98.14600186 98.50083692 98.84219798 99.14377836
 99.429229   99.65767636 99.84322373 100.         ]
```

```
[ ]: print(explained_variance)
```

```
[0.047045840868900575, 0.04353935201222255, 0.03952726132573601,
0.038998947946944876, 0.03754875514786483, 0.03450982750431003,
0.03384355853166483, 0.03209314365545552, 0.0315532966667674,
0.03013296958404461, 0.029078299182827613, 0.026825270020589445,
0.026137767889050735, 0.02567816584302076, 0.02508751295984726,
0.024269198595429805, 0.023615072576010455, 0.020953791964544076,
0.020031176619161, 0.019968236503696862, 0.01931850234868134,
0.018617782237923954, 0.01841660438951509, 0.015866362888087338,
0.015519164996488926, 0.015145255048816478, 0.014113867956005655,
0.013611659113954329, 0.013263055897155467, 0.012206223287504397,
0.011795083087986788, 0.01166488047421333, 0.011426946662708805,
0.01115343813489865, 0.010840743364865162, 0.01064903588829331,
0.010173082164622748, 0.009909385035477069, 0.009635676070379623,
0.009556962407438073, 0.009294901532381526, 0.008328807391888968,
0.007875661309277027, 0.007288631711811241, 0.0070862486286809405,
0.006890617590802156, 0.006786065419673662, 0.006694247940971714,
0.006243620736290022, 0.006163582918789114, 0.006003961685779088,
0.0056412975598951155, 0.005315335966438774, 0.005153507382075013,
0.0049213645119450195, 0.0045240338400346, 0.003926975591687665,
0.0035483506728697313, 0.0034136105161505117, 0.0030158038150721656,
0.0028545064578439873, 0.0022844735378660347, 0.0018554737128098162,
0.0015677627158603381]
```

```
[ ]: explained_variance_ratio=[]
for i in range(64):
    e=explained_variance[i]/np.sum(explained_variance)
    explained_variance_ratio.append(e)
print(explained_variance_ratio)
```

```
[0.047045840868900575, 0.04353935201222255, 0.03952726132573601,
```

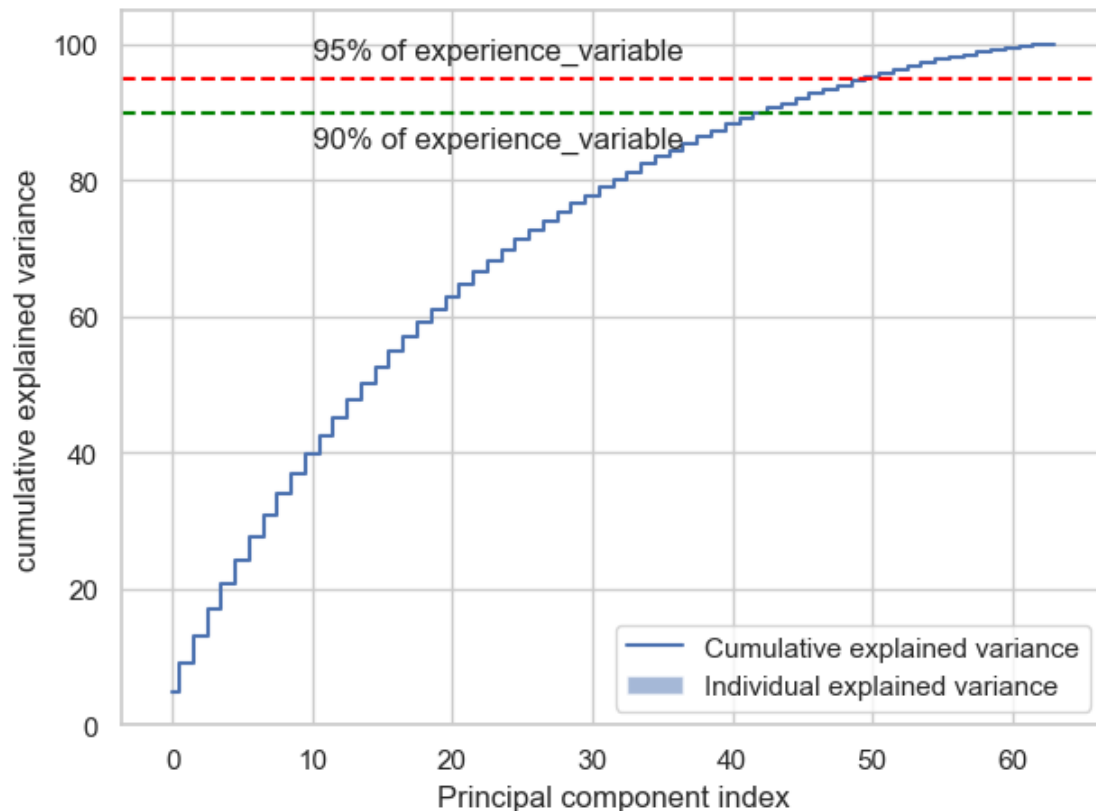
```
0.038998947946944876, 0.03754875514786483, 0.03450982750431003,
0.03384355853166483, 0.03209314365545552, 0.0315532966667674,
0.03013296958404461, 0.029078299182827613, 0.026825270020589445,
0.026137767889050735, 0.02567816584302076, 0.02508751295984726,
0.024269198595429805, 0.023615072576010455, 0.020953791964544076,
0.020031176619161, 0.019968236503696862, 0.01931850234868134,
0.018617782237923954, 0.01841660438951509, 0.015866362888087338,
0.015519164996488926, 0.015145255048816478, 0.014113867956005655,
0.013611659113954329, 0.013263055897155467, 0.012206223287504397,
0.011795083087986788, 0.01166488047421333, 0.011426946662708805,
0.01115343813489865, 0.010840743364865162, 0.01064903588829331,
0.010173082164622748, 0.009909385035477069, 0.009635676070379623,
0.009556962407438073, 0.009294901532381526, 0.008328807391888968,
0.007875661309277027, 0.007288631711811241, 0.0070862486286809405,
0.006890617590802156, 0.006786065419673662, 0.006694247940971714,
0.006243620736290022, 0.006163582918789114, 0.006003961685779088,
0.0056412975598951155, 0.005315335966438774, 0.005153507382075013,
0.0049213645119450195, 0.0045240338400346, 0.003926975591687665,
0.0035483506728697313, 0.0034136105161505117, 0.0030158038150721656,
0.0028545064578439873, 0.0022844735378660347, 0.0018554737128098162,
0.0015677627158603381]
```

```
[ ]: cumulative_variance_ratio_df = pd.DataFrame(cumulative_variance_ratio)

len(cumulative_variance_ratio_df.loc[cumulative_variance_ratio_df[0] <= 91])
```

```
[ ]: 44
```

```
[ ]: plt.bar(range(0,len(explained_variance)), explained_variance, alpha=0.5,
            ↪align='center', label='Individual explained variance')
plt.step(range(0,len(cumulative_variance_ratio)), cumulative_variance_ratio,
        ↪where='mid',label='Cumulative explained variance')
plt.ylabel('Explained variance ratio')
plt.xlabel('Principal component index')
plt.legend(loc='best')
plt.axhline(y=95,c='red',linestyle='--')
plt.axhline(y=90,c='green',linestyle='--')
plt.ylabel('cumulative explained variance')
plt.annotate('95% of experience_variable',
            xy=(10, 98),fontsize=12)
plt.annotate('90% of experience_variable',
            xy=(10, 85),fontsize=12)
plt.tight_layout()
plt.show()
```

KPCA with RBF Kernel

```
[ ]: from sklearn.decomposition import KernelPCA
import seaborn as sns
kpca = KernelPCA(kernel='rbf', n_components=64)
kpca=kpca.fit_transform(X_scaled_train)
explained_variance=np.var(kpca, axis=0 )

[ ]: explained_variance_ratio=explained_variance/ np.sum(explained_variance)*100
cumulative_variance_ratio=np.cumsum(explained_variance_ratio)
kpca_df=pd.DataFrame()
kpca_df['cumulative explained variance']= cumulative_variance_ratio
kpca_df['explained_variance_ratio']= explained_variance_ratio
display(kpca_df)
```

	cumulative explained variance	explained_variance_ratio
0	14.494873	14.494873
1	19.410752	4.915879
2	24.105188	4.694437
3	28.331324	4.226136
4	31.666812	3.335488
..


59	98.591645	0.385132
60	98.975189	0.383543
61	99.336178	0.360990
62	99.671726	0.335548
63	100.000000	0.328274

[64 rows x 2 columns]

```
[ ]: cumulative_variance_ratio_df = pd.DataFrame(cumulative_variance_ratio)
     len(cumulative_variance_ratio_df.loc[cumulative_variance_ratio_df[0] <= 91])
```

```
[ ]: 64
```

```
[ ]: plt.bar(range(0,len(explained_variance_ratio)), explained_variance_ratio,
             ↪alpha=0.5, align='center', label='Individual explained variance')
plt.step(range(0,len(cumulative_variance_ratio)), cumulative_variance_ratio,
         ↪where='mid',label='Cumulative explained variance')
plt.xlabel('Principal component index')
plt.legend(loc='best')
plt.axhline(y=91,c='black',linestyle='--')
plt.ylabel('cumulative explained variance')
plt.annotate('91% of experience_variable',
             xy=(0, 86),fontsize=12)
plt.tight_layout()
plt.show()
```




```
[ ]: from sklearn.decomposition import KernelPCA
Kernel_pca = KernelPCA(n_components = 44 ,kernel= "rbf",n_jobs=-1)# extracts 2
    ↳features, specify the kernel as rbf
X_train_kpca = Kernel_pca.fit_transform(X_scaled_train)
X_test_kpca= Kernel_pca.transform(X_scaled_test)

X_train_kpca=pd.DataFrame(X_train_kpca)
X_test_kpca=pd.DataFrame(X_test_kpca)
```

```
[ ]: linear_svm = SVC(kernel='linear')
linear_svm.fit(X_train_kpca, y_train)
y_pred_test = linear_svm.predict(X_test_kpca)
y_pred_train = linear_svm.predict(X_train_kpca)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
    ↳y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
    ↳y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.4443936971911395


Test accuracy: 0.4383561643835616



```
[ ]: rbf_svm = SVC(kernel='rbf')
rbf_svm.fit(X_train_kpca, y_train)
y_pred_test = rbf_svm.predict(X_test_kpca)
y_pred_train = rbf_svm.predict(X_train_kpca)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.9462205983101165

Test accuracy: 0.8873287671232877



```
[ ]: param_grid = [ {
    'C' : [4,6,10,20,25],
    "gamma" : ['scale',0.01,0.1,0.0001],
    'kernel' : ['rbf']},
    ]

model = SVC(kernel="rbf")

optimal_parameters =GridSearchCV(
    model, param_grid,
    cv=5,
    scoring= 'accuracy',
    verbose=10,
    n_jobs= -5,
    return_train_score=True
)
optimal_parameters.fit(X_train_kpca,y_train)


print(optimal_parameters.best_params_)
print(optimal_parameters.best_score_)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits
{'C': 4, 'gamma': 'scale', 'kernel': 'rbf'}
0.8917556569109888

```
[ ]: rbf_svm_better = SVC(kernel='rbf',C = 25, gamma = 'scale')
rbf_svm_better.fit(X_train_kpca, y_train)
y_pred_test = rbf_svm_better.predict(X_test_kpca)
y_pred_train = rbf_svm_better.predict(X_train_kpca)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.992235670244348


Test accuracy: 0.8993150684931507



```
[ ]: from sklearn.neighbors import KNeighborsClassifier
knn_model = KNeighborsClassifier(n_neighbors=2)
knn_model.fit(X_train_kpca, y_train)
y_pred_test = knn_model.predict(X_test_kpca)
y_pred_train = knn_model.predict(X_train_kpca)
print("Train accuracy: ", metrics.accuracy_score(y_true=y_train,
↪ y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪ y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.8619547841973053

Test accuracy: 0.751027397260274



```
[ ]: from sklearn.neighbors import NearestCentroid

ncc_model = NearestCentroid()
ncc_model.fit(X_train_kpca, y_train)
```


```

y_pred_test = ncc_model.predict(X_test_kpca)
y_pred_train = ncc_model.predict(X_train_kpca)
print("Train accuracy: ", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()

```

Train accuracy: 0.42943594427951587

Test accuracy: 0.4476027397260274



KPCA+LDA with RBF Kernel

```

[ ]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

lda = LDA(n_components=None)
X_train_kpca_lda = lda.fit_transform(X_train_kpca, y_train)

```



```
[ ]: # Set initial variance explained so far
total_variance = 0.0
# Set initial number of features
n_components = 0
# For the explained variance of each feature:
for explained_variance in lda.explained_variance_ratio_:
    # Add the explained variance to the total
    total_variance += explained_variance
    # Add one to the number of components
    n_components += 1

    # If we reach our goal level of explained variance
    if total_variance >= 0.91:
        break

# Return the number of components
print(n_components)
```


1

```
[ ]: lda = LDA(n_components=1)
X_train_kpca_lda = lda.fit_transform(X_train_kpca, y_train)
X_test_kpca_lda = lda.transform(X_test_kpca)
```

```
[ ]: linear_svm = SVC(kernel='linear')
linear_svm.fit(X_train_kpca_lda, y_train)
y_pred_test = linear_svm.predict(X_test_kpca_lda)
y_pred_train = linear_svm.predict(X_train_kpca_lda)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.4337748344370861


Test accuracy: 0.4452054794520548



```
[ ]: rbf_svm = SVC(kernel='rbf')
rbf_svm.fit(X_train_kpca_lda, y_train)
y_pred_test = rbf_svm.predict(X_test_kpca_lda)
y_pred_train = rbf_svm.predict(X_train_kpca_lda)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.38661794930349397

Test accuracy: 0.36198630136986304



```
[ ]: param_grid = [ {
    'C' : [4,6,10,20,25],
    "gamma" : ['scale',0.01,0.1,0.0001],
    'kernel' : ['rbf']},
    ]

model = SVC(kernel="rbf")

optimal_parameters =GridSearchCV(
    model, param_grid,
    cv=5,
    scoring= 'accuracy',
    verbose=10,
    n_jobs= -5,
    return_train_score=True
)
optimal_parameters.fit(X_train_kpca_lda,y_train)


print(optimal_parameters.best_params_)
print(optimal_parameters.best_score_)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits
{'C': 20, 'gamma': 'scale', 'kernel': 'rbf'}
0.4456489963986659

```
[ ]: rbf_svm_better = SVC(kernel='rbf',C=20,gamma='scale')
rbf_svm_better.fit(X_train_kpca_lda, y_train)
y_pred_test = rbf_svm_better.predict(X_test_kpca_lda)
y_pred_train = rbf_svm_better.predict(X_train_kpca_lda)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.4449646037908198


Test accuracy: 0.4606164383561644



```
[ ]: from sklearn.neighbors import KNeighborsClassifier
knn_model = KNeighborsClassifier(n_neighbors=2)
knn_model.fit(X_train_kpca_lda, y_train)
y_pred_test = knn_model.predict(X_test_kpca_lda)
y_pred_train = knn_model.predict(X_train_kpca_lda)
print("Train accuracy: ", metrics.accuracy_score(y_true=y_train,
↪ y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪ y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.689540991093857

Test accuracy: 0.3852739726027397



```
[ ]: from sklearn.neighbors import NearestCentroid

ncc_model = NearestCentroid()
ncc_model.fit(X_train_kpca_lda, y_train)
```


```

y_pred_test = ncc_model.predict(X_test_kpca_lda)
y_pred_train = ncc_model.predict(X_train_kpca_lda)
print("Train accuracy: ", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()

```

Train accuracy: 0.42098652660424757

Test accuracy: 0.44623287671232875



KPCA with Sigmoid Kernel

```

[ ]: from sklearn.decomposition import KernelPCA
import seaborn as sns
kpca = KernelPCA(kernel='sigmoid', n_components=64)
kpca=kpca.fit_transform(X_scaled_train)
explained_variance=np.var(kpca, axis=0 )

```

```
[ ]: explained_variance_ratio=explained_variance/ np.sum(explained_variance)
      cumulative_variance_ratio=np.cumsum(explained_variance_ratio)*100
      kpca_df=pd.DataFrame()
      kpca_df['cumulative explained variance']= cumulative_variance_ratio
      kpca_df['explained_variance_ratio']= explained_variance_ratio*100
      display(kpca_df)
```

	cumulative explained variance	explained_variance_ratio
0	4.779009	4.779009
1	9.190418	4.411409
2	13.171135	3.980717
3	17.086425	3.915291
4	20.864494	3.778069
..
59	99.157272	0.279456
60	99.429987	0.272715
61	99.648334	0.218347
62	99.830501	0.182167
63	100.000000	0.169499

[64 rows x 2 columns]

```
[ ]: cumulative_variance_ratio_df = pd.DataFrame(cumulative_variance_ratio)


      len(cumulative_variance_ratio_df.loc[cumulative_variance_ratio_df[0] <= 91])
```

[]: 44

```
[ ]: plt.bar(range(0,len(explained_variance_ratio)), explained_variance_ratio,α
      ↪alpha=0.5, align='center', label='Individual explained variance')
      plt.step(range(0,len(cumulative_variance_ratio)), cumulative_variance_ratio,α
      ↪where='mid',label='Cumulative explained variance')

      plt.xlabel('Principal component index')
      plt.legend(loc='best')
      plt.axhline(y=91,c='black',linestyle='--')
      plt.ylabel('cumulative explained variance')
      plt.annotate('91% of experience_variable',
                  xy=(1,86),fontsize=12)

      plt.tight_layout()
      plt.show()
```




```
[ ]: from sklearn.decomposition import KernelPCA
Kernel_pca = KernelPCA(n_components = 44 ,kernel= "sigmoid",n_jobs=-1)#
    ↳extracts 2 features, specify the kernel as rbf
X_train_kpca = Kernel_pca.fit_transform(X_scaled_train)
X_test_kpca= Kernel_pca.transform(X_scaled_test)

X_train_kpca=pd.DataFrame(X_train_kpca)
X_test_kpca=pd.DataFrame(X_test_kpca)
```

```
[ ]: linear_svm = SVC(kernel='linear')
linear_svm.fit(X_train_kpca, y_train)
y_pred_test = linear_svm.predict(X_test_kpca)
y_pred_train = linear_svm.predict(X_train_kpca)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
    ↳y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
    ↳y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.3034939483900434


Test accuracy: 0.27294520547945206



```
[ ]: rbf_svm = SVC(kernel='rbf')
rbf_svm.fit(X_train_kpca, y_train)
y_pred_test = rbf_svm.predict(X_test_kpca)
y_pred_train = rbf_svm.predict(X_train_kpca)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.9359442795158712

Test accuracy: 0.8900684931506849



```
[ ]: param_grid = [ {
    'C' : [4,6,10,20,25],
    "gamma" : ['scale',0.01,0.1,0.0001],
    'kernel' : ['rbf']},
    ]

model = SVC(kernel="rbf")

optimal_parameters =GridSearchCV(
    model, param_grid,
    cv=5,
    scoring= 'accuracy',
    verbose=2,
    n_jobs= -5,
    return_train_score=True
)
optimal_parameters.fit(X_train_kpca,y_train)


print(optimal_parameters.best_params_)
print(optimal_parameters.best_score_)
```

Fitting 5 folds for each of 20 candidates, totalling 100 fits
{'C': 6, 'gamma': 'scale', 'kernel': 'rbf'}
0.8941535691281434

```
[ ]: rbf_svm_better = SVC(kernel='rbf',C=6,gamma='scale')
rbf_svm_better.fit(X_train_kpca, y_train)
y_pred_test = rbf_svm_better.predict(X_test_kpca)
y_pred_train = rbf_svm_better.predict(X_train_kpca)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.9809317195706783


Test accuracy: 0.9102739726027397



```
[ ]: from sklearn.neighbors import KNeighborsClassifier
knn_model = KNeighborsClassifier(n_neighbors=2)
knn_model.fit(X_train_kpca, y_train)
y_pred_test = knn_model.predict(X_test_kpca)
y_pred_train = knn_model.predict(X_train_kpca)
print("Train accuracy: ", metrics.accuracy_score(y_true=y_train,
↪ y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪ y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.8622973281571135

Test accuracy: 0.6482876712328767



```
[ ]: from sklearn.neighbors import NearestCentroid

ncc_model = NearestCentroid()
ncc_model.fit(X_train_kpca, y_train)
```


```

y_pred_test = ncc_model.predict(X_test_kpca)
y_pred_train = ncc_model.predict(X_train_kpca)
print("Train accuracy: ", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()

```

Train accuracy: 0.29595798127426354

Test accuracy: 0.2660958904109589



```

[ ]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA

lda = LDA(n_components=None)
X_train_kpca_lda = lda.fit_transform(X_train_kpca, y_train)
# Set initial variance explained so far
total_variance = 0.0
# Set initial number of features

```

```

n_components = 0
# For the explained variance of each feature:
for explained_variance in lda.explained_variance_ratio_:
    # Add the explained variance to the total
    total_variance += explained_variance
    # Add one to the number of components
    n_components += 1

    # If we reach our goal level of explained variance
    if total_variance >= 0.95:
        break

# Return the number of components
print(n_components)

```

3

```

[ ]: lda = LDA(n_components=3)
X_train_kpca_lda = lda.fit_transform(X_train_kpca, y_train)
X_test_kpca_lda=lda.transform(X_test_kpca)

```


```

[ ]: linear_svm = SVC(kernel='linear')
linear_svm.fit(X_train_kpca_lda, y_train)
y_pred_test = linear_svm.predict(X_test_kpca_lda)
y_pred_train = linear_svm.predict(X_train_kpca_lda)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()

```

Train accuracy: 0.3142269924640329


Test accuracy: 0.27842465753424656



```
[ ]: rbf_svm = SVC(kernel='rbf')
rbf_svm.fit(X_train_kpca_lda, y_train)
y_pred_test = rbf_svm.predict(X_test_kpca_lda)
y_pred_train = rbf_svm.predict(X_train_kpca_lda)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪ y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪ y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.38661794930349397

Test accuracy: 0.36198630136986304



```
[ ]: param_grid = [ {
    'C' : [4,6,10,20,25],
    "gamma" : ['scale',0.01,0.1,0.0001],
    'kernel' : ['rbf']},
    ]

model = SVC(kernel="rbf")

optimal_parameters =GridSearchCV(
    model, param_grid,
    cv=5,
    scoring= 'accuracy',
    verbose=10,
    n_jobs= -5,
    return_train_score=True
)
optimal_parameters.fit(X_train_kpca_lda,y_train)


print(optimal_parameters.best_params_)
print(optimal_parameters.best_score_)
```


Fitting 5 folds for each of 20 candidates, totalling 100 fits
{'C': 4, 'gamma': 0.1, 'kernel': 'rbf'}
0.36777816459739904

```
[ ]: rbf_svm = SVC(kernel='rbf',C= 4, gamma= 0.1)
rbf_svm.fit(X_train_kpca_lda, y_train)
y_pred_test = rbf_svm.predict(X_test_kpca_lda)
y_pred_train = rbf_svm.predict(X_train_kpca_lda)
print("Train accuracy:", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.37645581182918475


Test accuracy: 0.3578767123287671



```
[ ]: from sklearn.neighbors import KNeighborsClassifier
knn_model = KNeighborsClassifier(n_neighbors=2)
knn_model.fit(X_train_kpca_lda, y_train)
y_pred_test = knn_model.predict(X_test_kpca_lda)
y_pred_train = knn_model.predict(X_train_kpca_lda)
print("Train accuracy: ", metrics.accuracy_score(y_true=y_train,
↪ y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪ y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()
```

Train accuracy: 0.6493491664763644

Test accuracy: 0.2839041095890411



```
[ ]: from sklearn.neighbors import NearestCentroid

ncc_model = NearestCentroid()
ncc_model.fit(X_train_kpca_lda, y_train)
```

```

y_pred_test = ncc_model.predict(X_test_kpca_lda)
y_pred_train = ncc_model.predict(X_train_kpca_lda)
print("Train accuracy: ", metrics.accuracy_score(y_true=y_train,
↪y_pred=y_pred_train), "\n")
print("Test accuracy:", metrics.accuracy_score(y_true=y_test,
↪y_pred=y_pred_test), "\n")
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_test)
plt.show()

```

Train accuracy: 0.3039506736697876

Test accuracy: 0.2684931506849315

